Interleaver Design for Deep Neural Networks

Sourya Dey, Peter Beerel, Keith Chugg

Asilomar Conference on Signals, Systems, and Computers

Oct-Nov 2017

USC University of Southern California

Overview of Current DNNs

- Key machine learning technologies

Overview of Current DNNs

- Key machine learning technologies

- Lot of parameters - **Memory intensive**
- Slow to train - **Computationally intensive**

Overview of Current DNNs

- Key machine learning technologies

- Lot of parameters - **Memory intensive**
- Slow to train - **Computationally intensive**

- Training done **offline** in CPU/GPU
- Custom hardware used for **inference only**

Typical Supervised Network

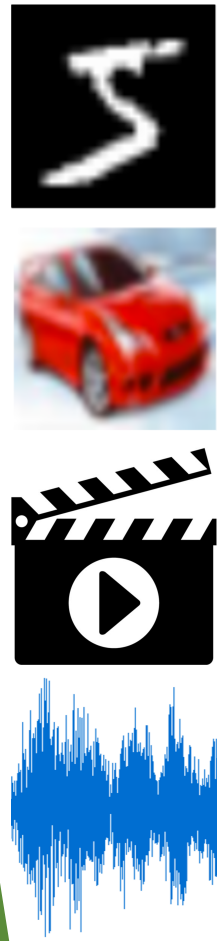

Convolutional Layers → Pooling and Normalization → Fully Connected Classification Layers (CLs) → **Predictions**

5% of total network connections

95% of total network connections

Krizhevsky, A., Sutskever, I., Hinton, G.E.: Imagenet classification with deep convolutional neural networks. In: NIPS-2012, pp. 1097–1105 (2012)
Zhang, C., Wu, D., Sun, J., Sun, G., Luo, G., Cong, J.: Energy-efficient CNN implementation on a deeply pipelined FPGA cluster. In: ISLPED-2016. pp. 326– 331. ACM, New York (2016)

Focus of our Approach

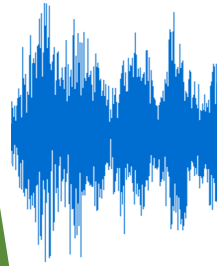

Convolutional Layers → Pooling and Normalization → **Fully Connected CLs** → **Predictions**

95% of total network connections

Overview of our Research

- Predefined sparsity - **Memory friendly**
  - *2-3x savings on CL only network parameters*
  - *2 orders of magnitude savings on CL parameters of CNNs **

Dey, S., Shao, Y., Chugg, K.M., Beerel, P.A.: Accelerating Training of Deep Neural Networks via Sparse Edge Processing. In: Proc. ICANN-2017, pp. 273-280. LNCS (2017)
* Dey, S., Huang, K.W., Beerel, P.A., Chugg, K.M.: Characterizing Sparse Connectivity Patterns in Neural Networks. In: ICLR-2018 (submitted for publication)

Overview of our Research

- Predefined sparsity - **Memory friendly**
  - *2-3x savings on CL only network parameters*
  - *2 orders of magnitude savings on CL parameters of CNNs **
- Edge-based processing - **Computationally flexible**
- Hardware optimizations - **Fast** training

Dey, S., Shao, Y., Chugg, K.M., Beerel, P.A.: Accelerating Training of Deep Neural Networks via Sparse Edge Processing. In: Proc. ICANN-2017, pp. 273-280. LNCS (2017)
* Dey, S., Huang, K.W., Beerel, P.A., Chugg, K.M.: Characterizing Sparse Connectivity Patterns in Neural Networks. In: ICLR-2018 (submitted for publication)

Overview of our Research

- Predefined sparsity - **Memory friendly**
  - *2-3x savings on CL only network parameters*
  - *2 orders of magnitude savings on CL parameters of CNNs **
- Edge-based processing - **Computationally flexible**
- Hardware optimizations - **Fast** training
- FPGA based architecture - **Online training** and inference

Dey, S., Shao, Y., Chugg, K.M., Beerel, P.A.: Accelerating Training of Deep Neural Networks via Sparse Edge Processing. In: Proc. ICANN-2017, pp. 273-280. LNCS (2017)
* Dey, S., Huang, K.W., Beerel, P.A., Chugg, K.M.: Characterizing Sparse Connectivity Patterns in Neural Networks. In: ICLR-2018 (submitted for publication)

Sparsity

Sparsity

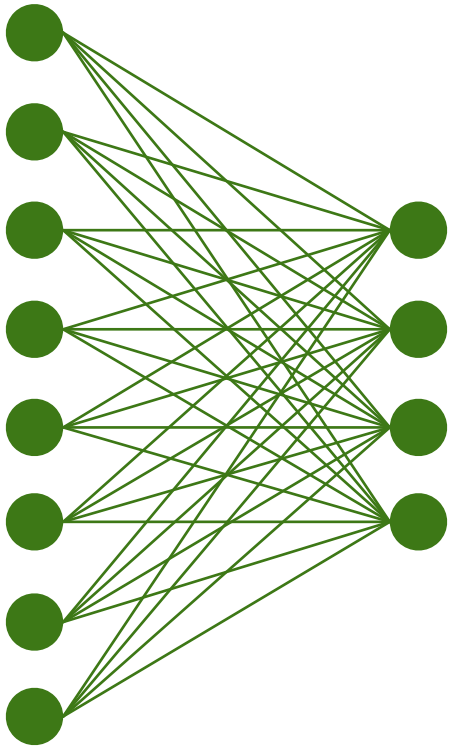

Fully connected (FC) network

Sparsity

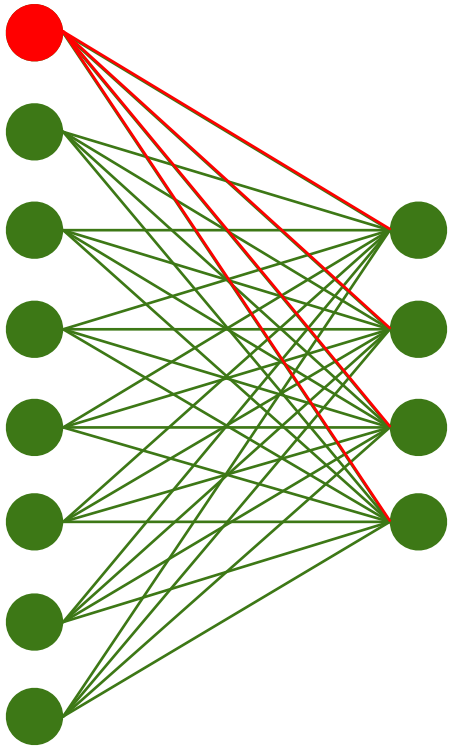

Fully connected (FC) network

Fanout (*fo*) = 4

Sparsity

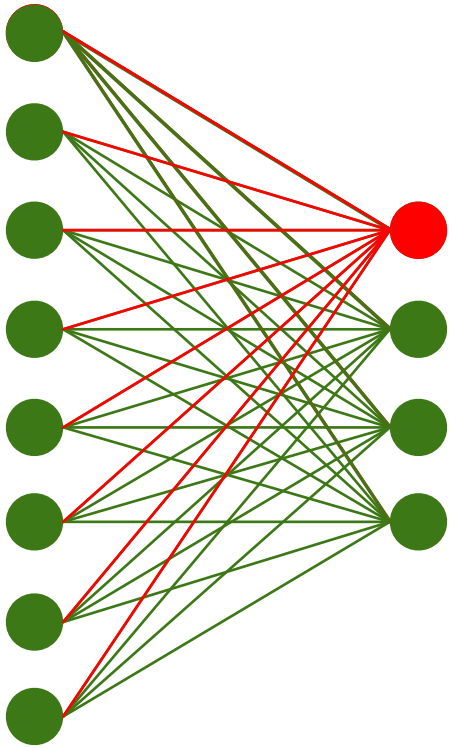

Fully connected (FC) network

Fanout (*fo*) = 4 Fanin (*fi*) = 8

Sparsity

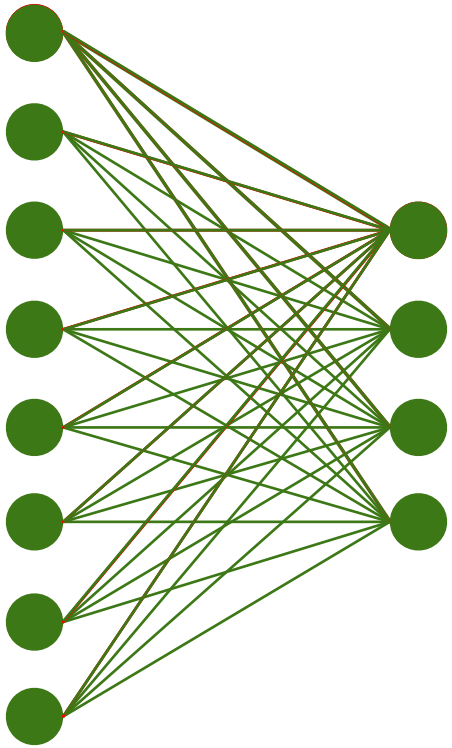

Fully connected (FC) network

Fanout (*fo*) = 4 Fanin (*fi*) = 8

Connectivity = 100%

Sparsity

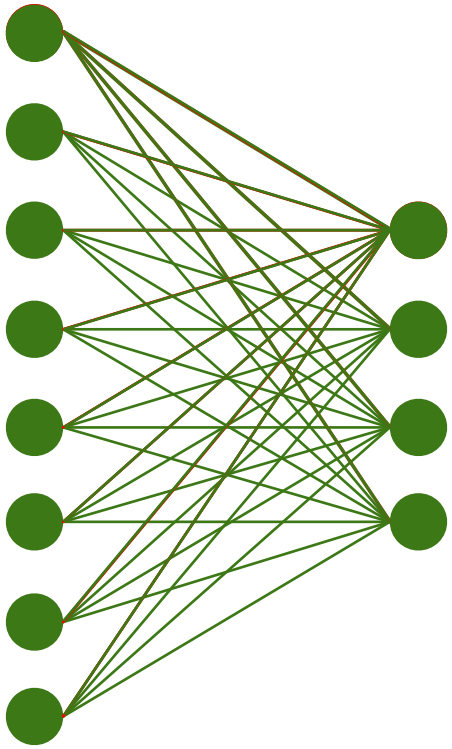

Fully connected (FC) network
Fanout (*fo*) = 4 Fanin (*fi*) = 8
Connectivity = 100%

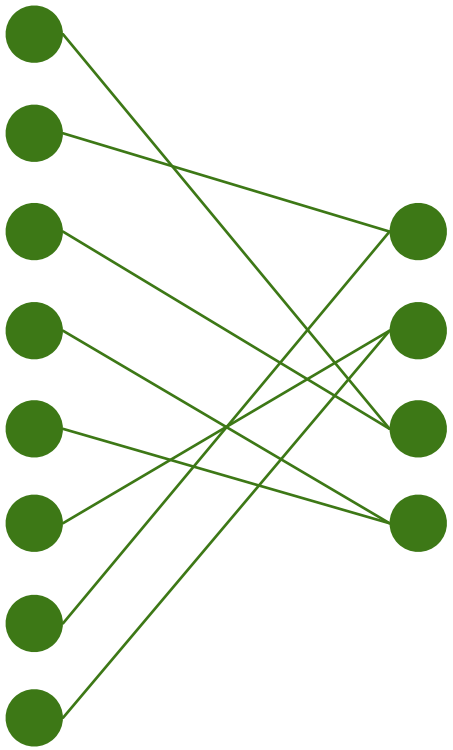

Sparse network
fo = 1, *fi* = 2
Connectivity = 25%

Sparsity – Predefined

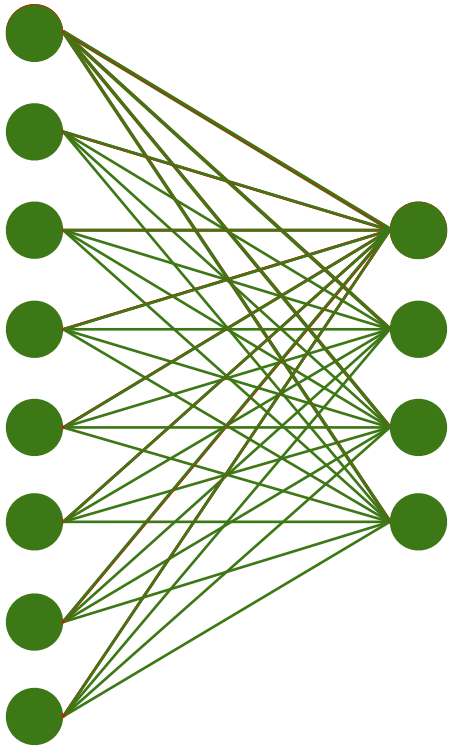

Fully connected (FC) network
Fanout (*fo*) = 4 Fanin (*fi*) = 8
Connectivity = 100%

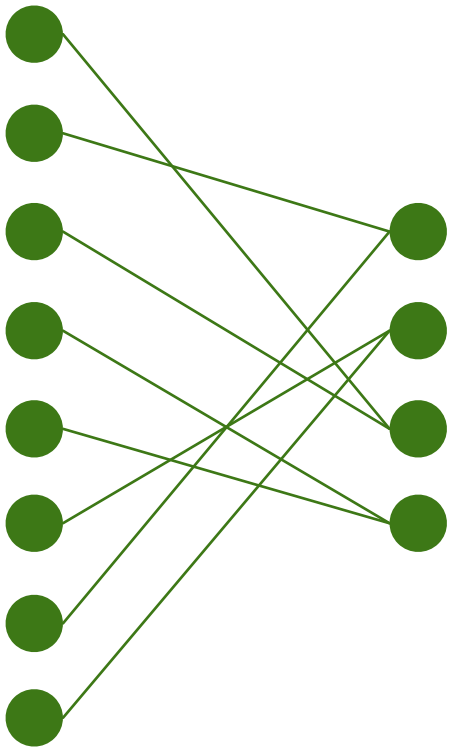

Sparse network
fo = 1, *fi* = 2
Connectivity = 25%

Example of Parameter Savings

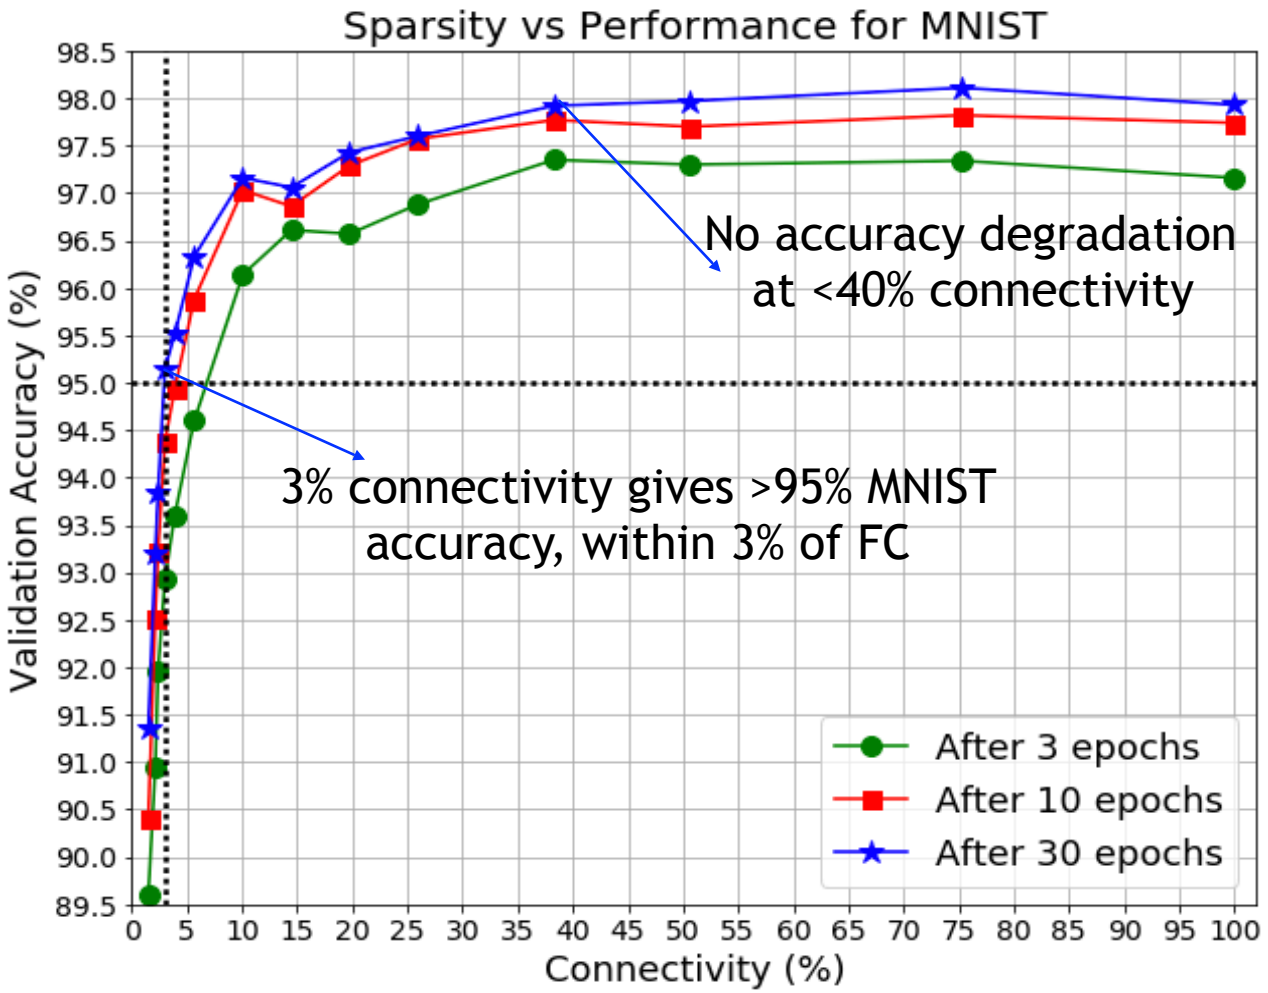

Present Work - Interleavers for Sparse Patterns

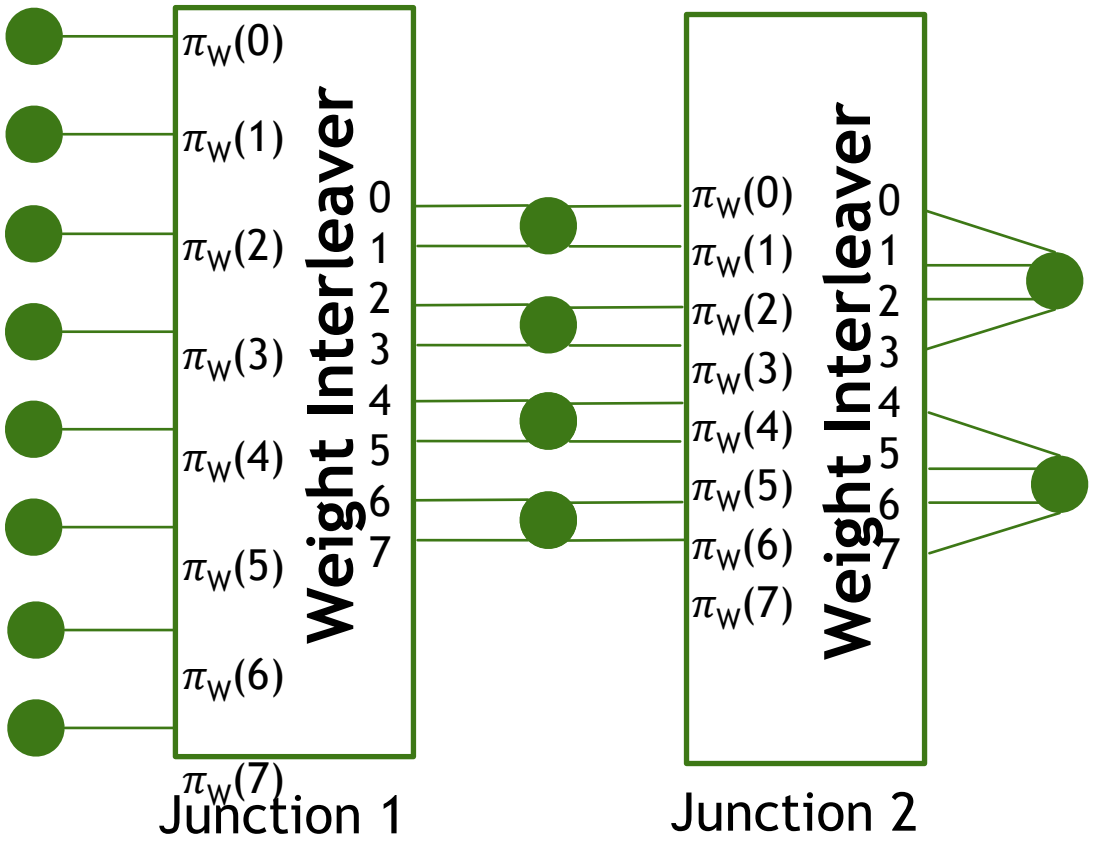

Present Work - Interleavers for Sparse Patterns

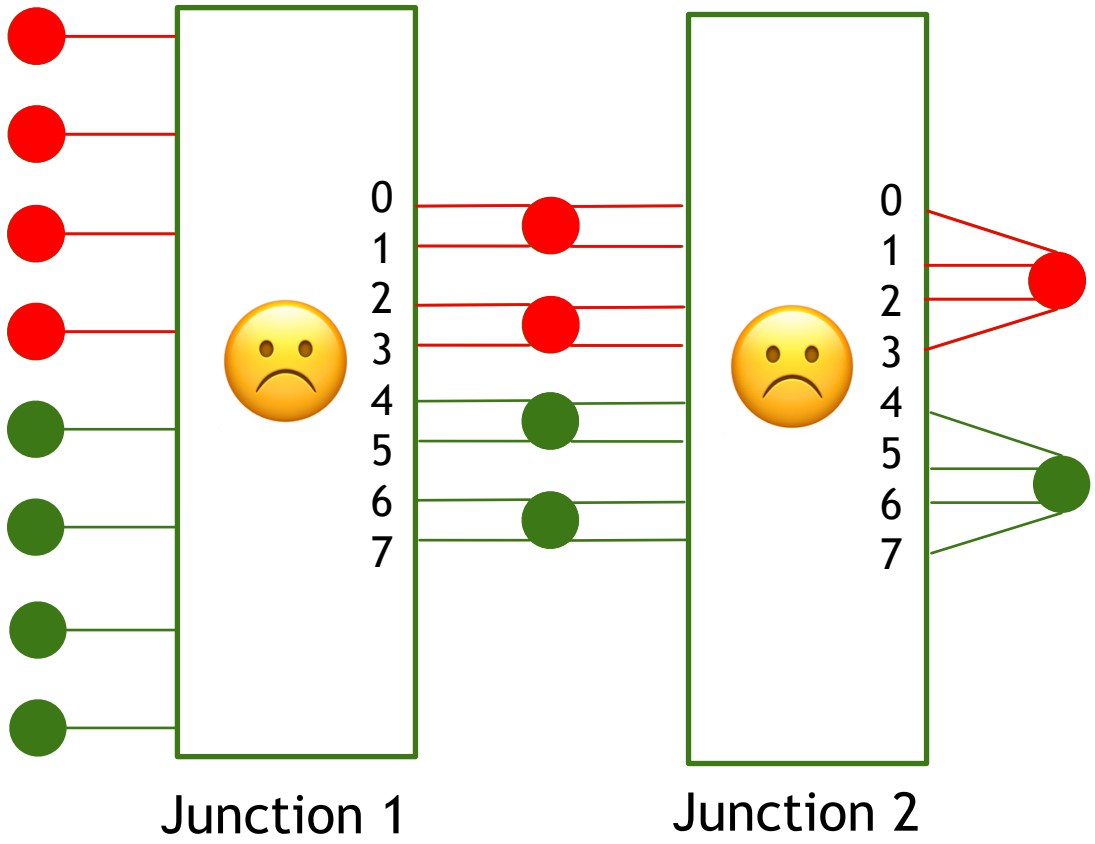

Present Work - Interleavers for Sparse Patterns

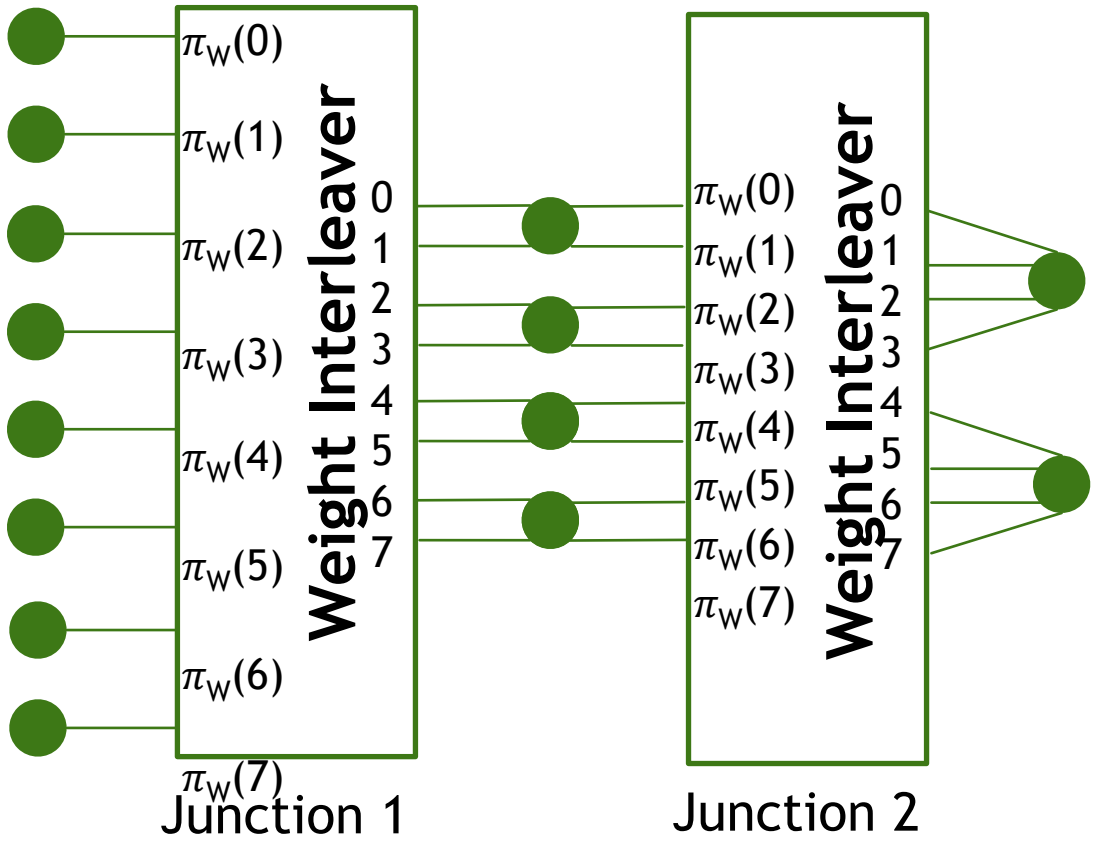

Interleaver algorithm ensures:

- Each output connected to a *good spatial chunk* of different inputs
- No neuron unconnected

Interleaver Requirements

- Optimized for computational efficiency in hardware
- Optimized for on-chip storage
- High values for metrics which are performance indicators

Degree of Parallelism = z

Degree of Parallelism = z

- z memories for all parameters of same type

Mem 0	Mem 1				Mem z-1
W_0	W_1	W_2			W_{z-1}
W_z					
W_{2z}					
W_{3z}					W_{W-1}

Degree of Parallelism = z

- z memories for all parameters of same type
- Process z parameters in 1 **cycle** => 1 from each mem

Mem 0	Mem 1				Mem z-1
W_0	W_1	W_2			W_{z-1}
W_z					
W_{2z}					
W_{3z}					W_{W-1}

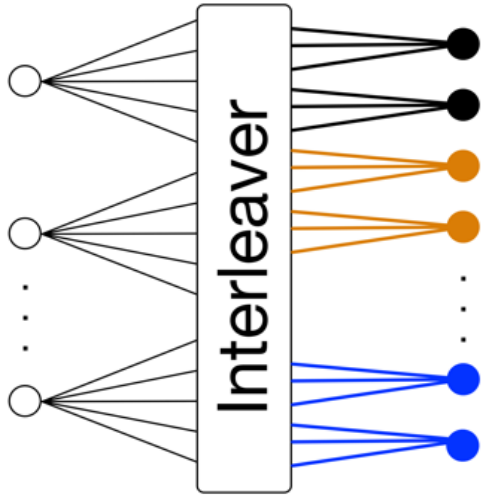

Degree of Parallelism = z

- z memories for all parameters of same type
- Process z parameters in 1 **cycle** => 1 from each mem
- **Process all parameters in a sweep**

Mem 0	Mem 1				Mem z-1
W_0	W_1	W_2			W_{z-1}
W_z					
W_{2z}					
W_{3z}					W_{W-1}

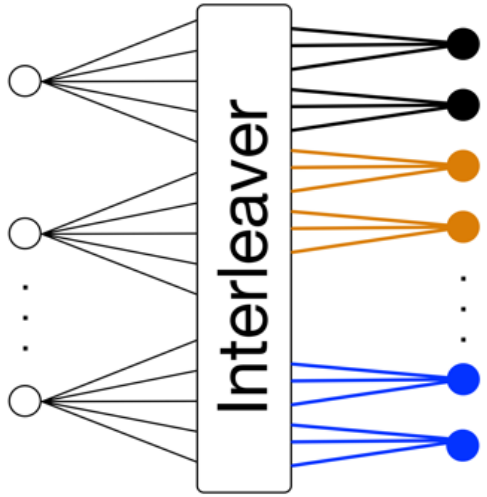

Clash Freedom

Clash Freedom

- E.g. p activations, so depth of each memory = p/z
- Accessed in interleaved (permuted order)

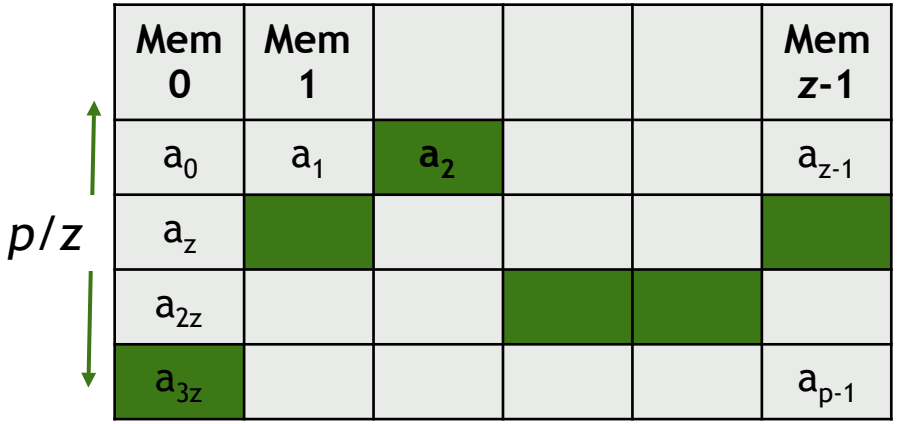

Clash Freedom

- E.g. p activations, so depth of each memory = p/z
- Accessed in interleaved (permuted order)

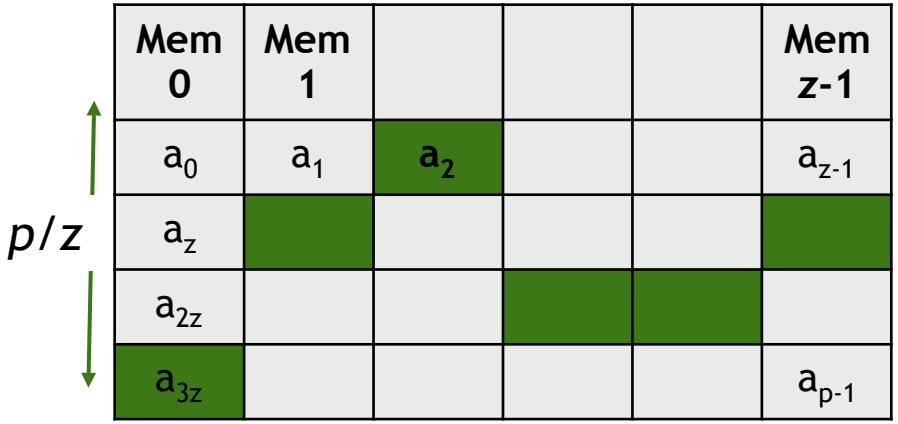

Clash-free access ☺️

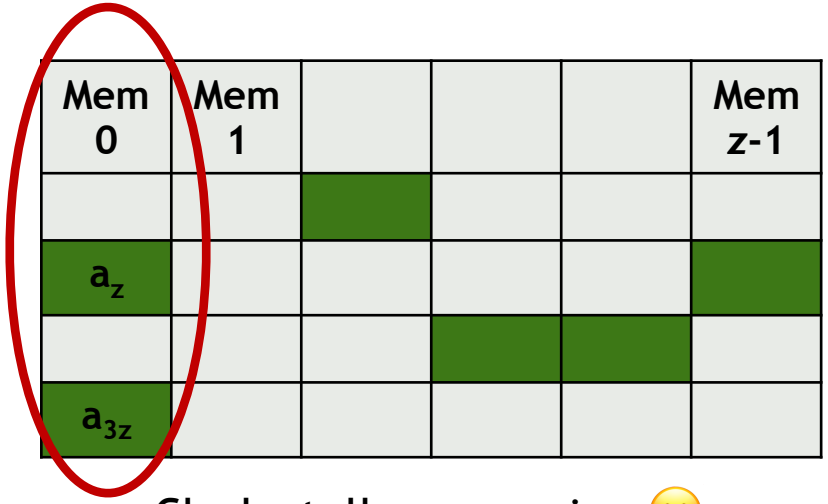

Clash stalls processing ☹️

Clash Freedom

- E.g. p activations, so depth of each memory = p/z
- Accessed in interleaved (permuted order)

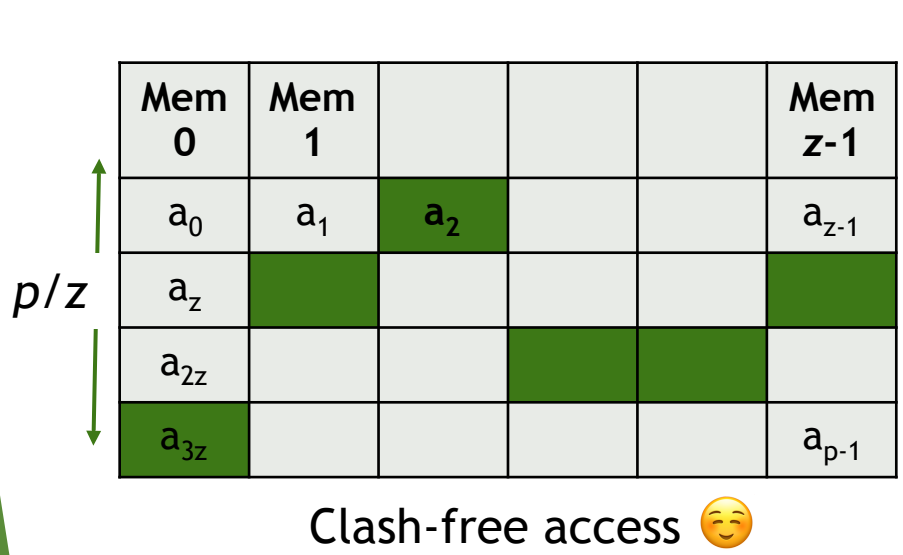

Clash-free access ☺

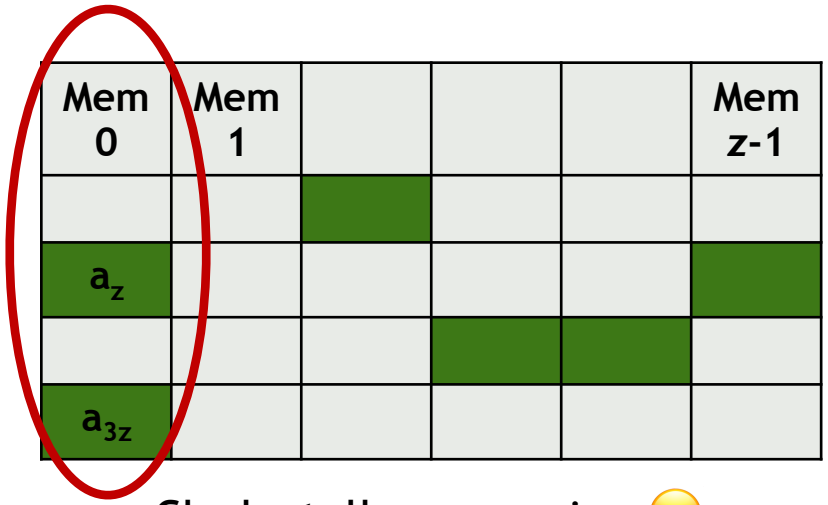

Clash stalls processing ☹

Interleaver must prevent clashes when accessing activations

Ease of Accesses

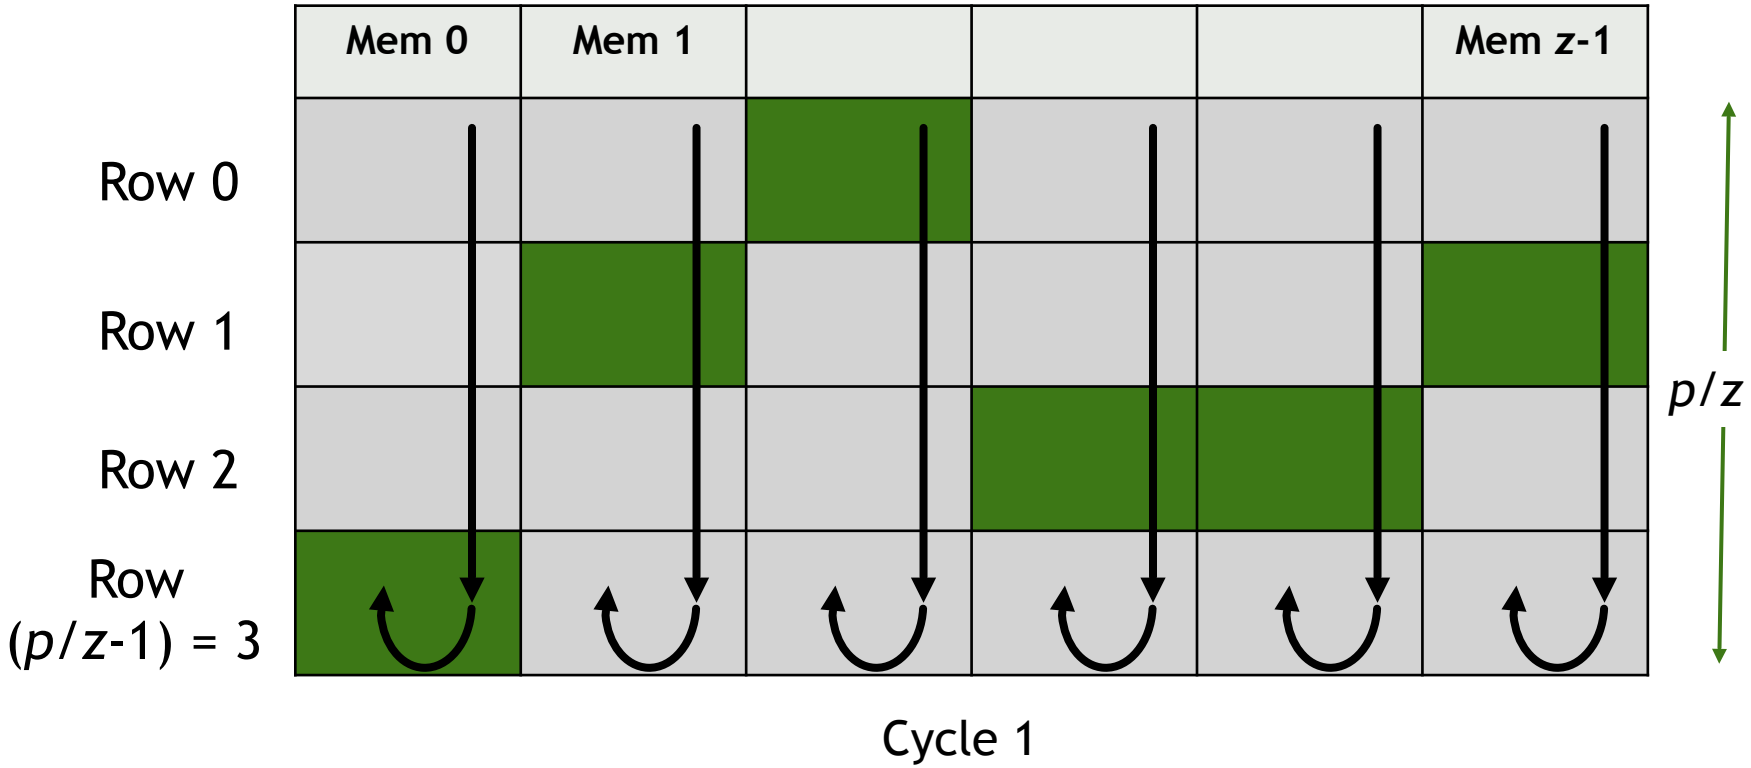

Cycle 1

Design interleaver to have easy address computation

Ease of Accesses

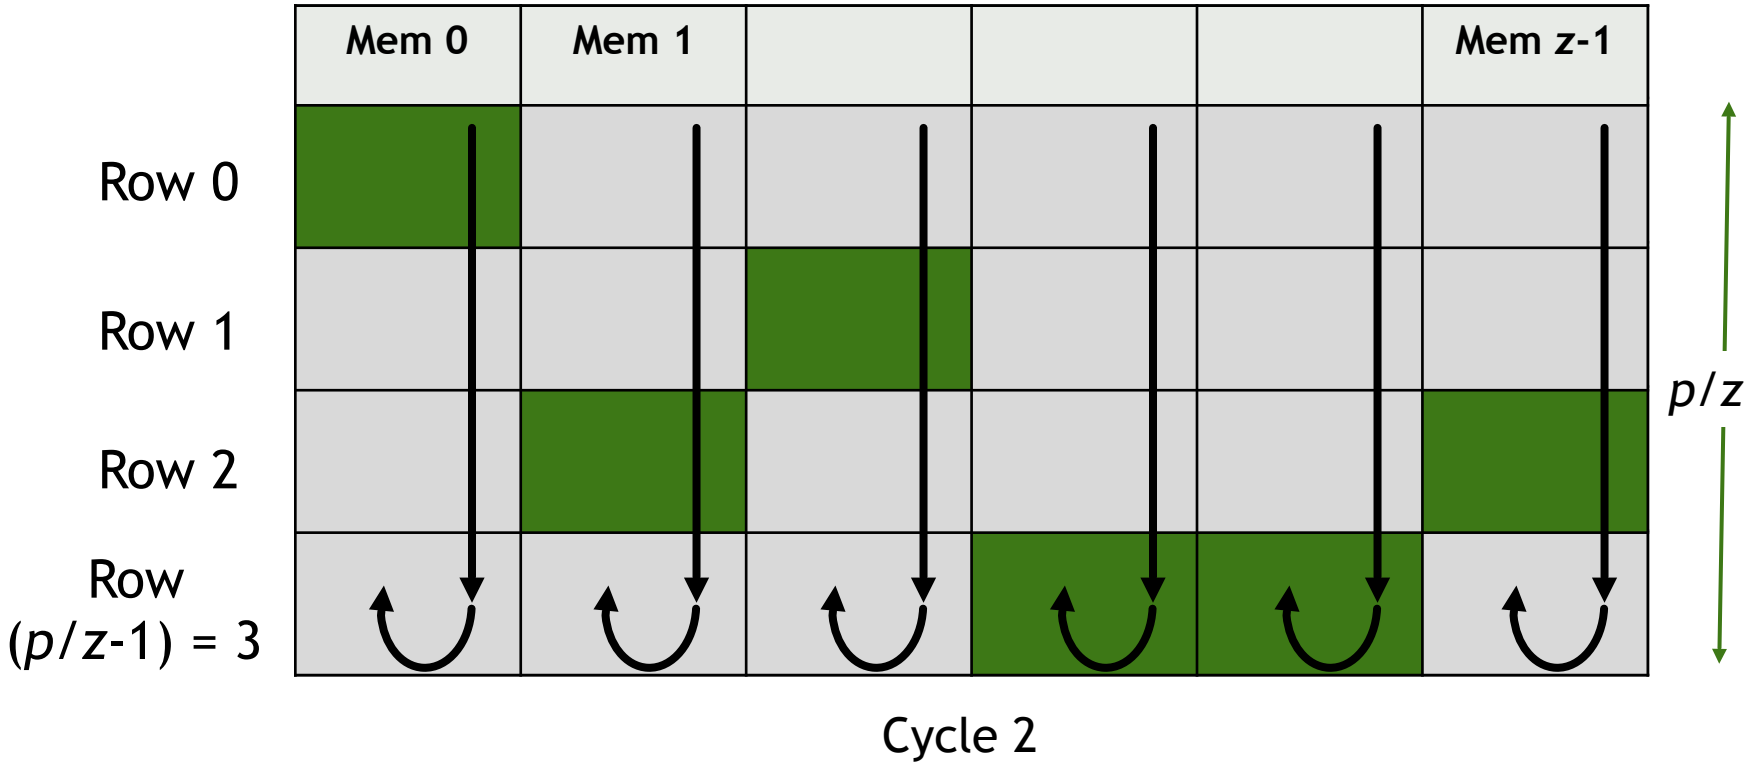

Cycle 2

Design interleaver to have easy address computation

Ease of Accesses

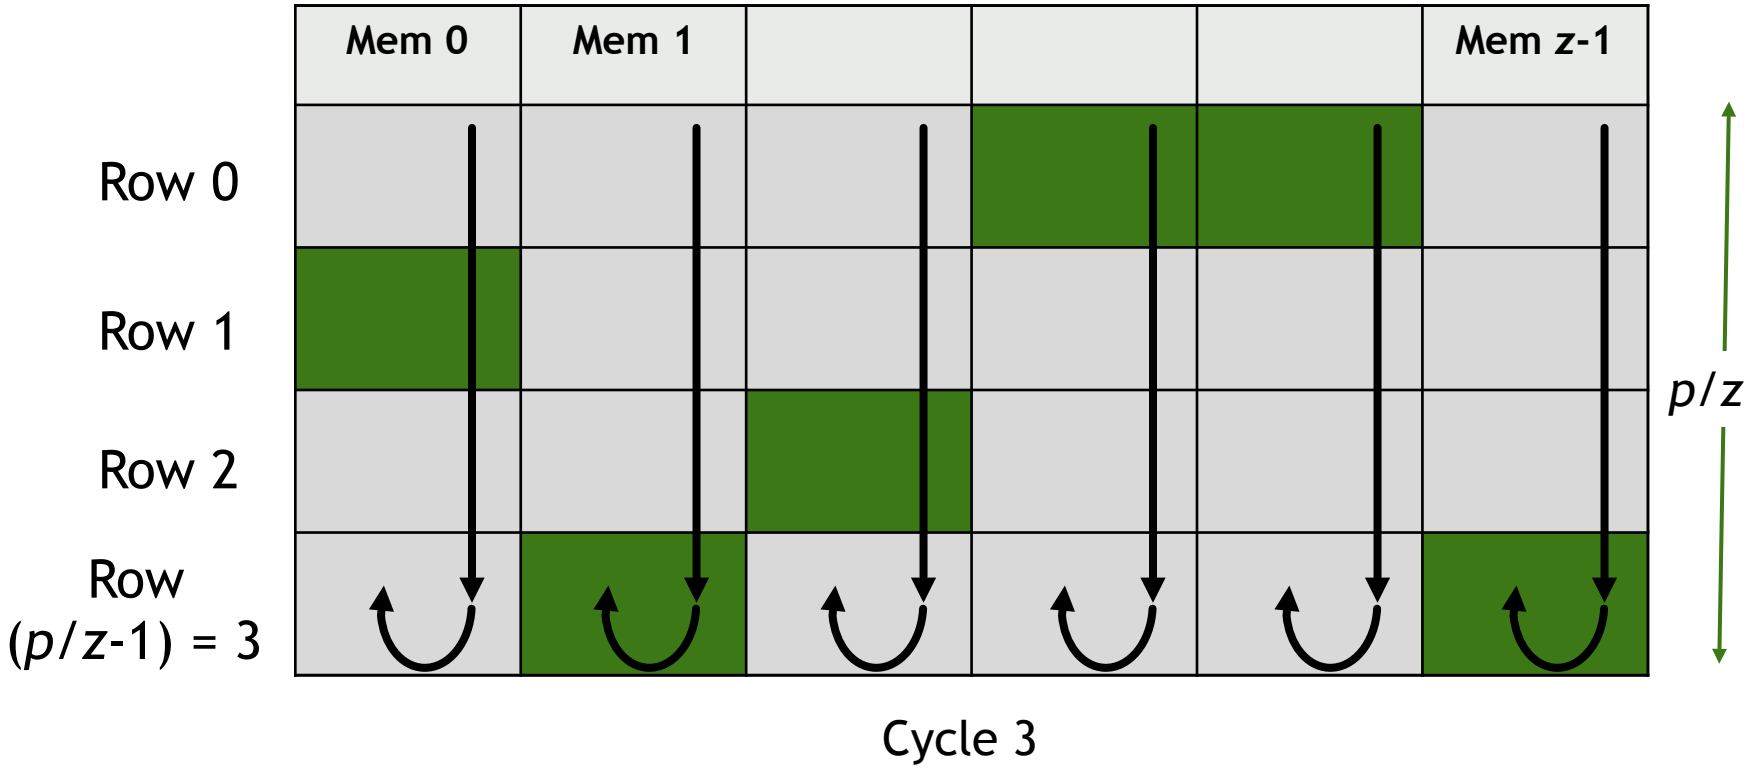

Cycle 3

Design interleaver to have easy address computation

Ease of Accesses

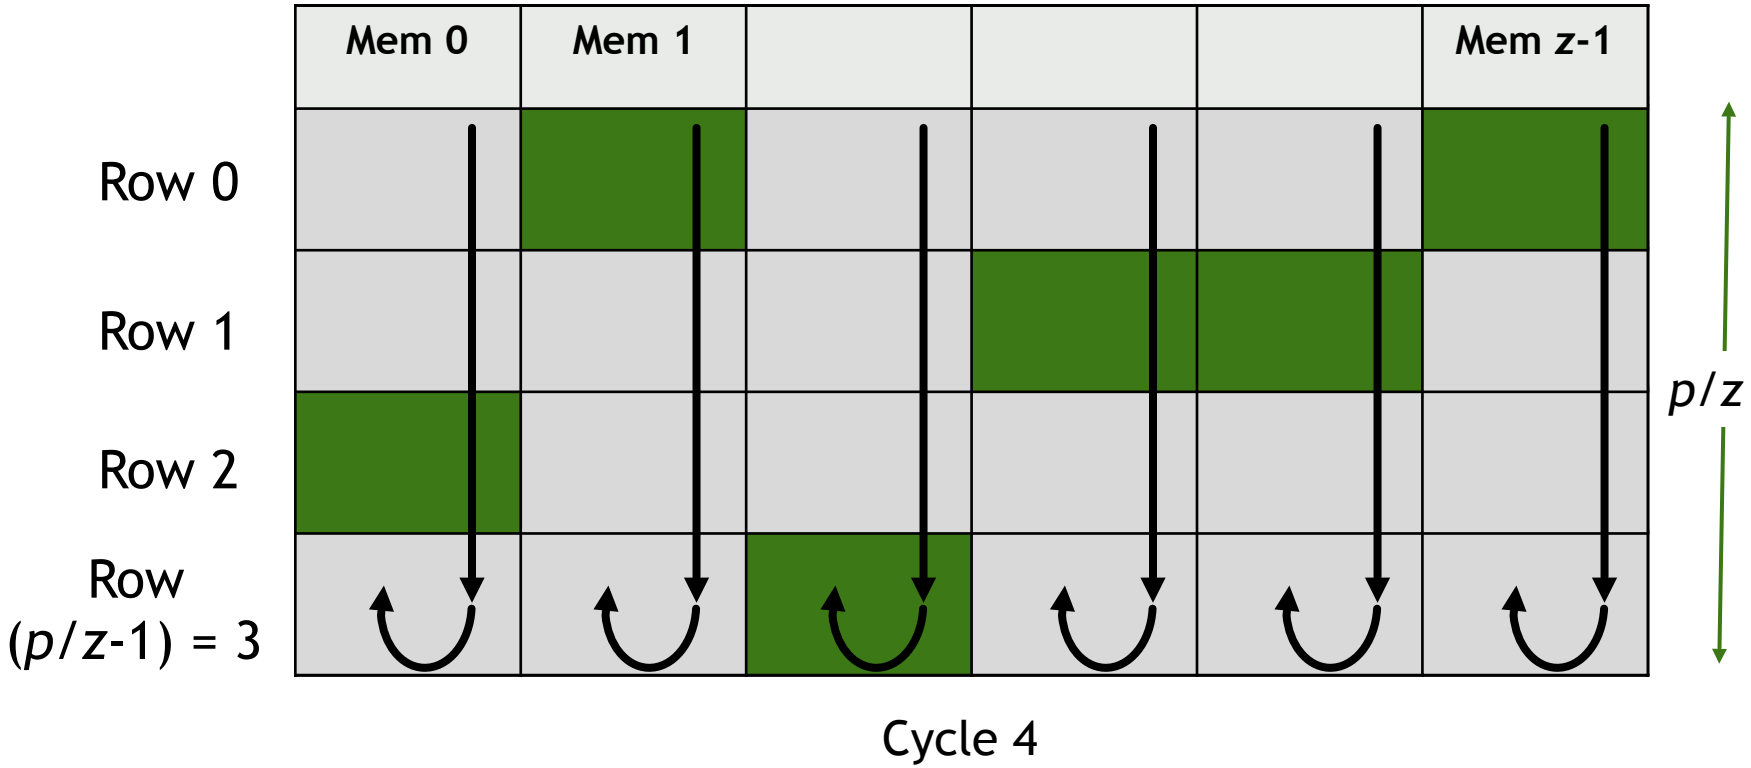

Cycle 4

Design interleaver to have easy address computation

Interleaver Design Algorithm

Interleaver Design Algorithm

- Let ***r*** be a random permutation of memory row index => Size p/z

Interleaver Design Algorithm

- Let ***r*** be a random permutation of memory row index => Size *p*/*z*
- Replicate or partition *r* to form ***s*** of *z* elements => Starting indices of all mems

Interleaver Design Algorithm

- Let ***r*** be a random permutation of memory row index => Size *p*/*z*
- Replicate or partition *r* to form ***s*** of *z* elements => Starting indices of all mems
- ***t*** = {*s*, *s*+1, …, *s*+*p*/*z*-1}%(p/z) => All *p* indices for all mems in order

Interleaver Design Algorithm

- Let ***r*** be a random permutation of memory row index => Size *p*/*z*
- Replicate or partition *r* to form ***s*** of *z* elements => Starting indices of all mems
- ***t*** = {*s*, *s*+1, …, *s*+*p*/*z*-1}%(p/z) => All *p* indices for all mems in order

$$\pi_W(i) = \left(t\left[i\%p\right] \times z + i\%z \right) \times fo + \lfloor i/p \rfloor$$

Interleaver Design Algorithm

- Let ***r*** be a random permutation of memory row index => Size *p*/*z*
- Replicate or partition *r* to form ***s*** of *z* elements => Starting indices of all mems
- ***t*** = {*s*, *s*+1, …, *s*+*p*/*z*-1}%(p/z) => All *p* indices for all mems in order

Example: *p*=32, *fo*=2, *z*=8 => *i*∈[0,63]. Say *i* = 45

$$\pi_W(i) = \left(t\left[i\%p\right] \times z + i\%z\right) \times fo + \lfloor i/p \rfloor$$

Interleaver Design Algorithm

- Let ***r*** be a random permutation of memory row index => Size *p*/*z*
- Replicate or partition *r* to form ***s*** of *z* elements => Starting indices of all mems
- ***t*** = {*s*, *s*+1, …, *s*+*p*/*z*-1}%(p/z) => All *p* indices for all mems in order

Example: *p*=32, *fo*=2, *z*=8 => *i*∈[0,63]. Say *i* = 45

$$\pi_W(i) = \left(t\left[i\%p\right] \times z + i\%z \right) \times fo + \lfloor i/p \rfloor$$

Activation Memory Bank Row = 45%32 = 1

Row1	a_8	a_9	a_{10}	a_{11}	a_{12}	a_{13}	a_{14}	a_{15}

Interleaver Design Algorithm

- Let ***r*** be a random permutation of memory row index => Size *p*/*z*
- Replicate or partition *r* to form ***s*** of *z* elements => Starting indices of all mems
- ***t*** = {*s*, *s*+1, …, *s*+*p*/*z*-1}%(p/z) => All *p* indices for all mems in order

Example: *p*=32, *fo*=2, *z*=8 => *i*∈[0,63]. Say *i* = 45

$$\pi_W(i) = \left(t\left[i\%p\right] \times z + i\%z \right) \times fo + \lfloor i/p \rfloor$$

Activation Memory Bank Row = 45%32 = 1

Row1	a_8	a_9	a_{10}	a_{11}	a_{12}	a_{13}	a_{14}	a_{15}

a_{13} **Left side Neuron = 1x8+45%8 = 13**

Interleaver Design Algorithm

- Let ***r*** be a random permutation of memory row index => Size *p*/*z*
- Replicate or partition *r* to form ***s*** of *z* elements => Starting indices of all mems
- ***t*** = {*s*, *s*+1, …, *s*+*p*/*z*-1}%(p/z) => All *p* indices for all mems in order

Example: *p*=32, *fo*=2, *z*=8 => *i*∈[0,63]. Say *i* = 45

$$\pi_W(i) = \left(t[i\%p] \times z + i\%z \right) \times fo + \lfloor i/p \rfloor$$

Activation Memory Bank Row = 45%32 = 1

Row1	a_8	a_9	a_{10}	a_{11}	a_{12}	a_{13}	a_{14}	a_{15}

a_{13} **Left side Neuron = 1x8+45%8 = 13**

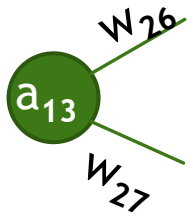

Left side Neuron's Weight = 13x2 = 26

Interleaver Design Algorithm

- Let ***r*** be a random permutation of memory row index => Size *p*/*z*
- Replicate or partition *r* to form ***s*** of *z* elements => Starting indices of all mems
- ***t*** = {*s*, *s*+1, …, *s*+*p*/*z*-1}%(p/z) => All *p* indices for all mems in order

Example: *p*=32, *fo*=2, *z*=8 => *i*∈[0,63]. Say *i* = 45

$$\pi_W(i) = \left(t[i\%p] \times z + i\%z \right) \times fo + \lfloor i/p \rfloor$$

Activation Memory Bank Row = 45%32 = 1

Row1	a_8	a_9	a_{10}	a_{11}	a_{12}	a_{13}	a_{14}	a_{15}

a_{13} **Left side Neuron = 1x8+45%8 = 13**

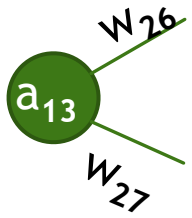

Left side Neuron's Weight = 13x2 = 26

Weight Offset = 1

Interleaver Design Algorithm

- Let ***r*** be a random permutation of memory row index => Size *p*/*z*
- Replicate or partition *r* to form ***s*** of *z* elements => Starting indices of all mems
- ***t*** = {*s*, *s*+1, …, *s*+*p*/*z*-1}%(p/z) => All *p* indices for all mems in order

Example: *p*=32, *fo*=2, *z*=8 => *i*∈[0,63]. Say *i* = 45

$$\pi_W(i) = \left(t[i\%p] \times z + i\%z \right) \times fo + \lfloor i/p \rfloor$$

Activation Memory Bank Row = 45%32 = 1

Row1	a_8	a_9	a_{10}	a_{11}	a_{12}	a_{13}	a_{14}	a_{15}

a_{13} **Left side Neuron = 1x8+45%8 = 13**

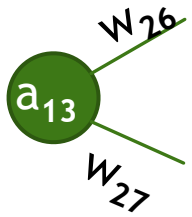

Left side Neuron's Weight = 13x2 = 26

Weight Offset = 1

Finally: w_{27}

Meeting Requirements

- Easily generated – Proof in paper
  - All variables involved are powers of 2 (add extra neurons)
    - Modulo = Bit select
    - Multiply = Bit shift
  - Only store *r* for a new pattern
    - Create t by accumulating 1s
- Clash freedom – Proof in paper

Variations

Variations

- Start Vector Shuffle (SV)

Original: s = {2,0,3,1,2,0,3,1} After SV: s = {2,0,3,1,**3,0,1,2**}

Variations

- Start Vector Shuffle (SV)

  Original: s = {2,0,3,1,2,0,3,1} After SV: s = {2,0,3,1,**3,0,1,2**}

- Sweep Starter Shuffle (SS)

  Original:
  1st sweep s = {2,0,3,1,2,0,3,1}
  2nd sweep s = {2,0,3,1,2,0,3,1}

  After SS:
  1st sweep s = {2,0,3,1,3,0,1,2}
  2nd sweep s = {**0,3,2,1,0,3,2,1**}

Variations

- Start Vector Shuffle (SV)

  Original: s = {2,0,3,1,2,0,3,1} After SV: s = {2,0,3,1,**3,0,1,2**}

- Sweep Starter Shuffle (SS)

  Original:
  1st sweep s = {2,0,3,1,2,0,3,1}
  2nd sweep s = {2,0,3,1,2,0,3,1}

  After SS:
  1st sweep s = {2,0,3,1,3,0,1,2}
  2nd sweep s = {**0,3,2,1,0,3,2,1**}

- Memory Dither (MD)

$$\pi_W(i) = \left(t\left[i\%p\right] \times z + \boldsymbol{v}[i\%z]\right) \times fo + \lfloor i/p \rfloor$$

v[.] = Permutation of [0,z-1]

Some Weight Interleaver Patterns

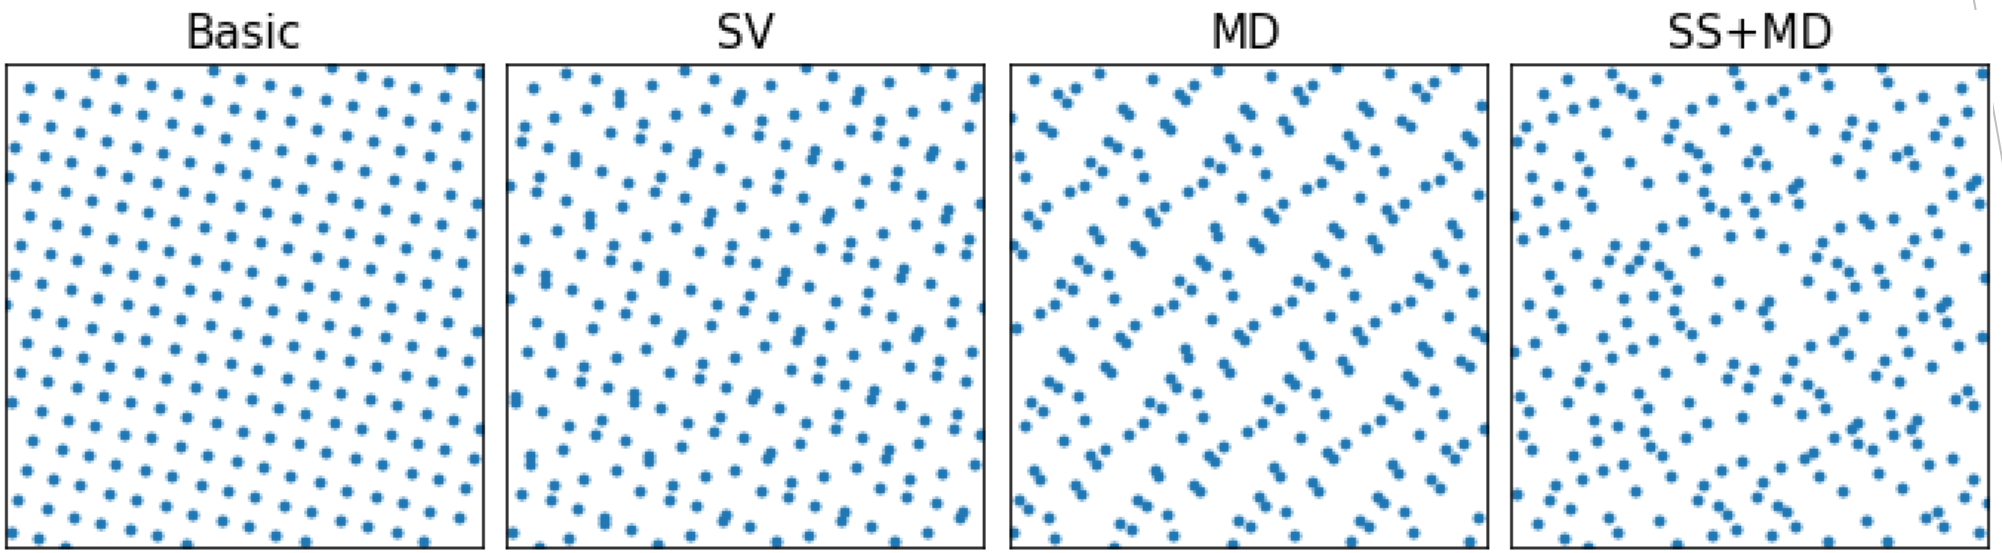

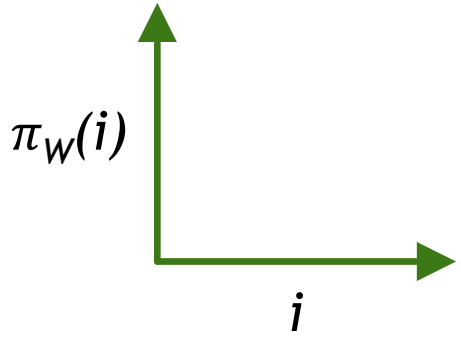

Spread and Dispersion

- Spread: Connections that are close on 1 side should be far away on other
- Dispersion: Connections should be irregular, i.e. no patterns or trends

Spread and Dispersion

- Spread: Connections that are close on 1 side should be far away on other
- Dispersion: Connections should be irregular, i.e. no patterns or trends

Interleaver Variant	Spread	Dispersion
Basic	**18.28**	0.04
MD	7.48	0.22
SS	9.7	0.07
SS + MD	6.5	0.37
SV	6.6	0.08
SV + MD	7.31	0.23
SV + SS	5.05	0.09
SV + SS + MD	5.7	**0.39**

Dataset Results

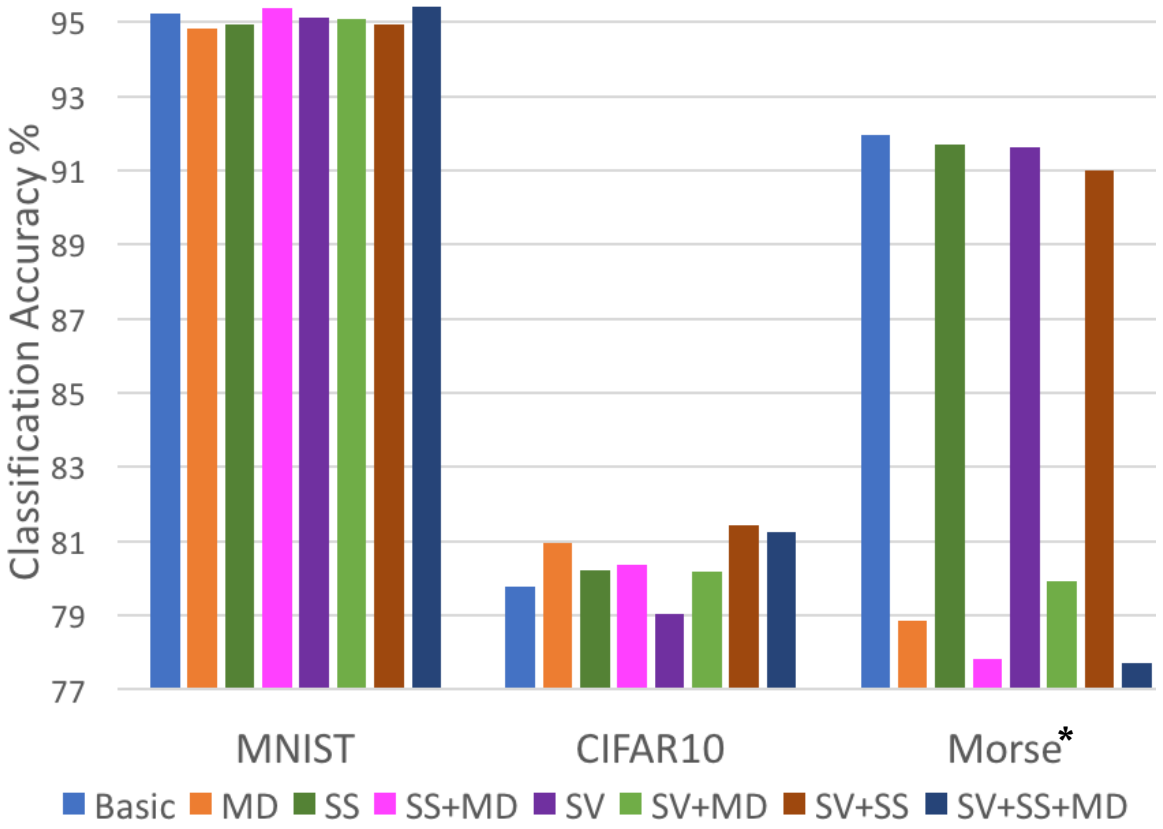

* Sourya Dey: https://cobaltfolly.wordpress.com/2017/10/15/morse-code-dataset-for-artificial-neural-networks/

Morse Dataset Trends

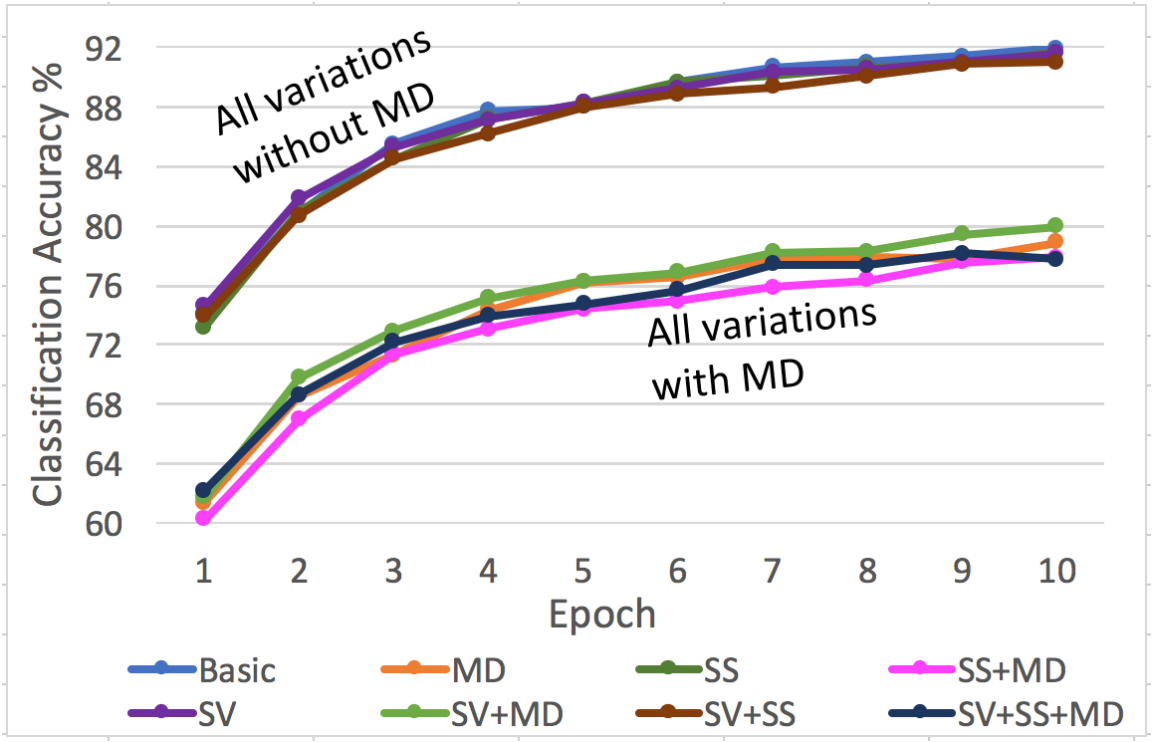

Morse has fewer inputs and low redundancy
Spread should be high, dispersion hurts

Summary and Ongoing Work

- Pre-defined sparse hardware architecture to lower memory and computational footprint
- Interleaver algorithm to guarantee clash freedom and ease of storage
- Interleaver variations and effects on performance

- Extension to multiple junctions - Adjacency matrices
- Measures to characterize network performance

Dey, S., Huang, K.W., Beerel, P.A., Chugg, K.M.: Characterizing Sparse Connectivity Patterns in Neural Networks. In: ICLR-2018 (submitted for publication)

Thank you!

Questions?

Contact: souryade@usc.edu